

Lab 09 - Golden test set to Latency budget

TGS-2024044051 | v1.0 | 27 September 2026

Scenario

Northstar Service is a synthetic service organisation. All records and metrics in this lab are fictional. Use an approved training tenant if available; otherwise complete the same evidence analysis offline.

Materials

- scenario.csv: 5 synthetic case records with task volumes, timing, claims and access values.
- roles.csv: pilot roles, licences and approved data access.
- pilot-metrics.csv: synthetic active seats, wage and license assumptions, and latency timings.
- ui-reference.png: authentic Tertiary Infotech training-tenant UI example; your tenant may differ.
- source-pack.md: synthetic policy excerpts, an outdated source and an injection test.
- evidence-template.md: learner evidence form.
- Microsoft 365 Copilot or Copilot Studio only when your trainer has provided access.

Procedure

1. Golden test set. Open the matching row in scenario.csv, then check its role against roles.csv and source against source-pack.md. Record the source version and owner. Version expected results and sources. Calculate net minutes saved as baseline minus pilot minus review, and verify supported claims do not exceed total claims. Save the output as 01-golden-test-set.md. Check: Data owner approves test updates. Negative test: Test set contains stale expected answer. Measure: Changed results per release.
2. Time-saved estimate. Open the matching row in scenario.csv, then check its role against roles.csv and source against source-pack.md. Record the source version and owner. Subtract review time and scale by volume. Calculate net minutes saved as baseline minus pilot minus review, and verify supported claims do not exceed total claims. Save the output as 02-time-saved-estimate.md. Check: Validate observed samples. Negative test: Gross saving ignores review effort. Measure: Net minutes saved per task.
3. ROI model. Open the matching row in scenario.csv, then check its role against roles.csv and source against source-pack.md. Record the source version and owner. Compare realized value to total costs. Calculate net minutes saved as baseline minus pilot minus review, and verify supported claims do not exceed total claims. Save the output as 03-roi-model.md. Check: Separate measured from assumed inputs. Negative test: Benefit double-counts the same task. Measure: (Value minus cost) divided by cost.
4. License optimization. Open the matching row in scenario.csv, then check its role against roles.csv and source against source-pack.md. Record the source version and owner. Reclaim idle seats after review. Calculate net minutes saved as baseline minus pilot minus review, and verify supported claims do not exceed total claims. Save the output as 04-license-optimization.md. Check: Line manager confirms exceptions. Negative test: Idle seats kept while pilot waits. Measure: Active seats divided by assigned seats.
5. Latency budget. Open the matching row in scenario.csv, then check its role against roles.csv and source against source-pack.md. Record the source version and owner. Decompose p50 and p95 response time. Calculate net minutes saved as baseline minus pilot minus review, and verify supported claims do not exceed total claims. Save the output as 05-latency-budget.md. Check: Do not remove safety checks for speed. Negative test: Tool call dominates p95. Measure: p95 seconds per task.

Copy-ready Copilot prompt

You are supporting a synthetic Northstar Service training case. Use only the provided scenario row. Identify your sources and assumptions. Complete the requested transformation, then list every claim that requires human verification. Never send, publish, grant access, or change production data.

Acceptance

- Five output artifacts correspond to the five rows in scenario.csv.
- Recompute net_minutes_saved independently from the three timing columns; record the source version used.
- For ROI use the role's hourly value and license cost from pilot-metrics.csv as labelled synthetic assumptions.
- For latency, add retrieval, model and tool component times for the same percentile; do not mix p50 with p95.
- Record whether the role may open the source; do not use the injected or retired source as an instruction.
- Every artifact records source, method, reviewer, date and a pass/fail decision.
- At least one failed or uncertain test is recorded with a correction.
- No real customer or tenant data appears in submitted evidence.

Troubleshooting

If the Microsoft feature is unavailable, state the licence/policy limitation and use the synthetic row to complete the same decision and validation evidence. Do not fabricate a live configuration.

UI reference

[!!!llustrative Microsoft workflow from a separate synthetic training tenant](#) This screenshot orients you to a Microsoft workspace. It is not evidence of your own configuration; submit your own screenshot or clearly labelled offline artifact.